

The Quest to Memory Safety

Programming Languages

Sebastian Fernandez
@snfernandez

Microsoft Security Response Center
Vulnerabilities & Mitigations



Who am I?

- Vulnerabilities and Mitigations team in MSRC
- Analyze and promote the use of safer languages
- Bug hunts and sometimes writes exploits

Defining the problem

Software vulnerabilities incur a high cost for everyone

For vendors

Enforcing secure development practices (SDL)

Fixing, verifying and shipping security patches

Bad PR

For customers

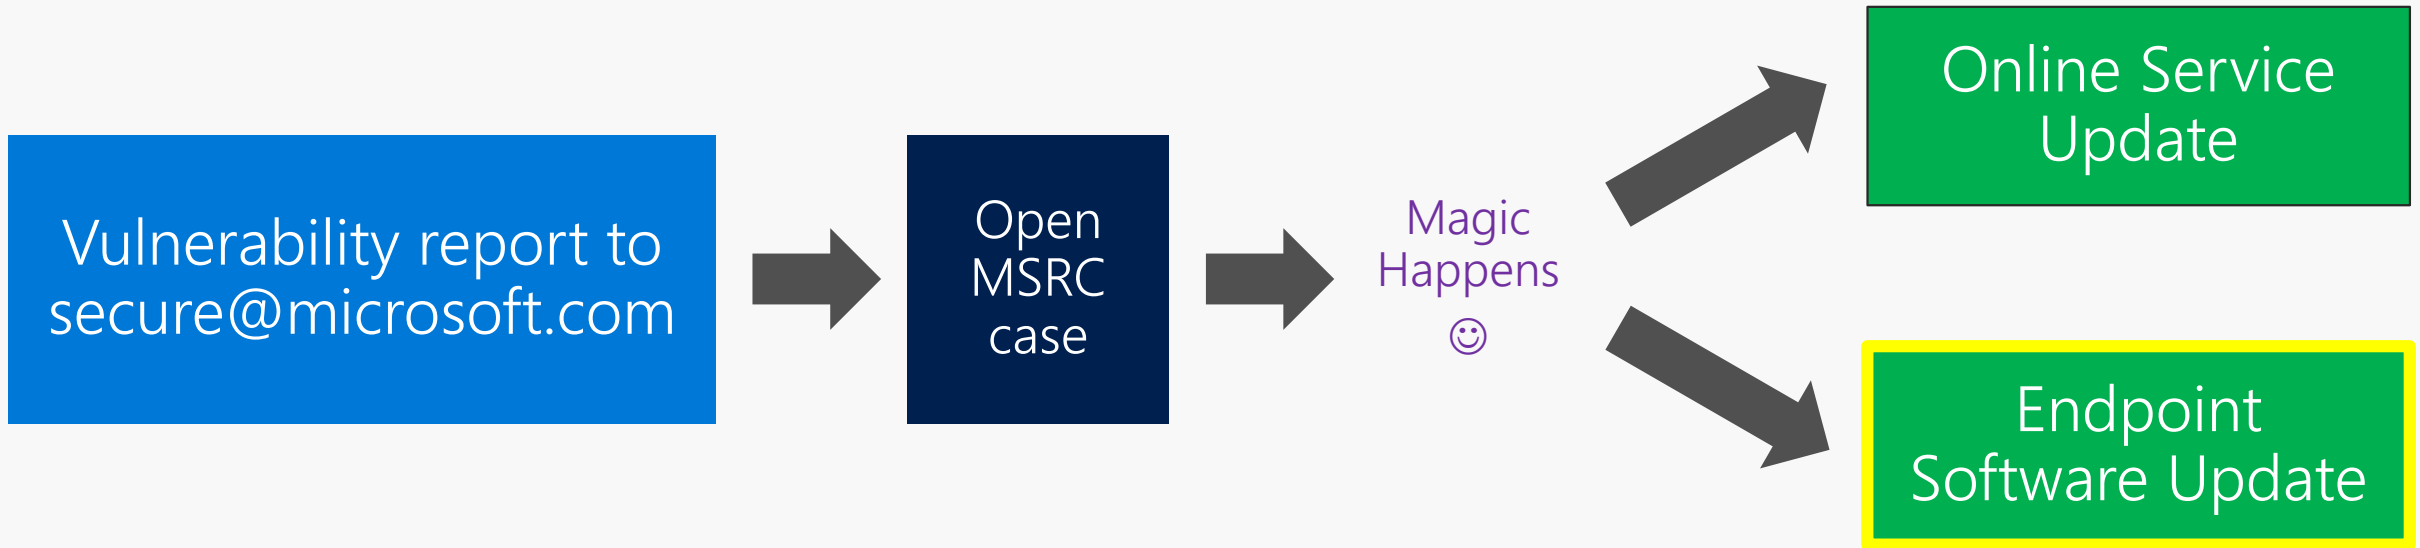
Vulnerable software leaves computers open to attacks

Extra layers of security

Disruptive security updates

Defining our scope

Vulnerabilities reported to Microsoft are typically addressed in one of two ways




In 2018, ~54% of reported vulnerabilities were addressed via a software update

~85% of those vulnerabilities were Remote Code Execution (RCE), Elevation of Privilege (EOP), or Information Disclosure (ID)


Memory safety issues remain dominant

We closely study the root cause trends of vulnerabilities & search for patterns



~70% of the vulnerabilities addressed through a security update each year continue to be memory safety issues

Drilling down into root causes



Top root causes since 2016:

#1: heap out-of-bounds

#2: use after free

#3: type confusion

#4: uninitialized use

And we are not alone

*"Most of Android's vulnerabilities occur in the media and Bluetooth components. **Use-after-free (UAF), integer overflows, and out of bounds (OOB) reads/writes comprise 90% of vulnerabilities** with OOB being the most common." – Jeff Vander Stoep about bugs in Android during 2018*

*"Across the entirety of iOS 12 Apple has fixed 261 CVEs, 173 of which were memory unsafety. **That's 66.3% of all vulnerabilities.** Across the entirety of Mojave Apple has fixed 298 CVEs, 213 of which were memory unsafety. **That's 71.5% of all vulnerabilities.**" – Paul Kehrler about bugs in Apple software*

*"There's a significant overlap between memory vulnerabilities and severe security problems. **Of the 34 critical/high bugs, 32 were memory-related.**" – Diane Hosfelt about bugs in the Firefox's styling code*

Microsoft's vulnerability mitigation strategy

Break exploitation techniques

- StackCookies, DEP, Heap Isolation, ACG, CFG, CIG

Eliminate vulnerabilities

- NULL deref protection, MemGC, InitAll

Contain damage & prevent persistence

- AppContainer, Virtualization

Promote safer programming languages

- .NET, Rust, C++17 & Core Guidelines

Transitioning to safer languages

```
string &not_fun(string &str, Ending *ending) {
    string str = str;
    for(auto it = str.begin(); it != vec.end(); ++it){
        auto v = *it;
        if (v >= 65) {
            str.push_back(v+12);
        } else {
            vec.erase(it);
        }
    }
    str.push_back(reinterpret_cast<AsciiEnding*>(ending)->get_char());
    return str;
}
```

Not a smart pointer

Might be a dangling pointer

Dangling pointer

Undefined behavior if signed char overflows

Invalidates iterator

Hmmm....

No NULL check

Returns a local dangling reference

Invalid cast

Systems programming

Systems programming aims to produce software and software platforms which provide services to other software, are performance constrained, or both

Desired features

Raw control over memory access and control flow

Easily interoperate with existing ABIs

Little runtime overhead

Software examples

Kernels and Hypervisors

Runtime libraries

Low latency (servers, game engines, browsers)

Which are our options?



C++ \geq 17
& Core Guidelines

- New features to help the user stay memory safe
- Easiest to integrate with existing C/C++ codebases
- No boundaries between unsafe and safe code

Not strong enough safety guarantees



Garbage Collected languages (C#, F#, TypeScript, Go, etc)

TS

- Memory is managed by the runtime
- Applies bounds checks at runtime
- Impact on performance predictability

Doesn't meet all "SPL" criteria



Rust

- Memory safe by default
- Fearless concurrency: data race safe by default
- Allows memory unsafety in a well-defined scope
- Performance comparable to C++

Memory safety

Spatial memory safety

Pointer indexing

Buffer overflows

Pointers lack size information

Time of check vs time of use

Arithmetic operations

Temporal memory safety

Use after free

Double free

No ownership model

Lack of lifetime enforcement

Type safety

Incorrect casting

Undefined behavior

Unsafe operations

Undefined states

Spatial Memory Safety

Spatial memory safety

Allocates
memory



Copies more than
allocated

```
char *ptr1;  
ptr1 = alloc(1)  
memset(mem, 0, 2)
```

Checked memory access

Transfers the responsibility of verifying the bounds of the objects to the runtime or enforcing them statically. Safeguards the user from making mistakes checking sizes

Fat pointers
(pointer+size)

- Rust `std::slice<T>`
- C++ `gsl::span<T>`

Arbitrary pointer
dereferencing

- No ☺. Well, you can still use *unsafe*

Temporal memory safety

Temporal memory safety



Enter memory ownership

A system that allows the developer to enforce when memory can be freed, transferred or moved, thus preventing them from accidentally creating and dereferencing dangling pointers

Single ownership

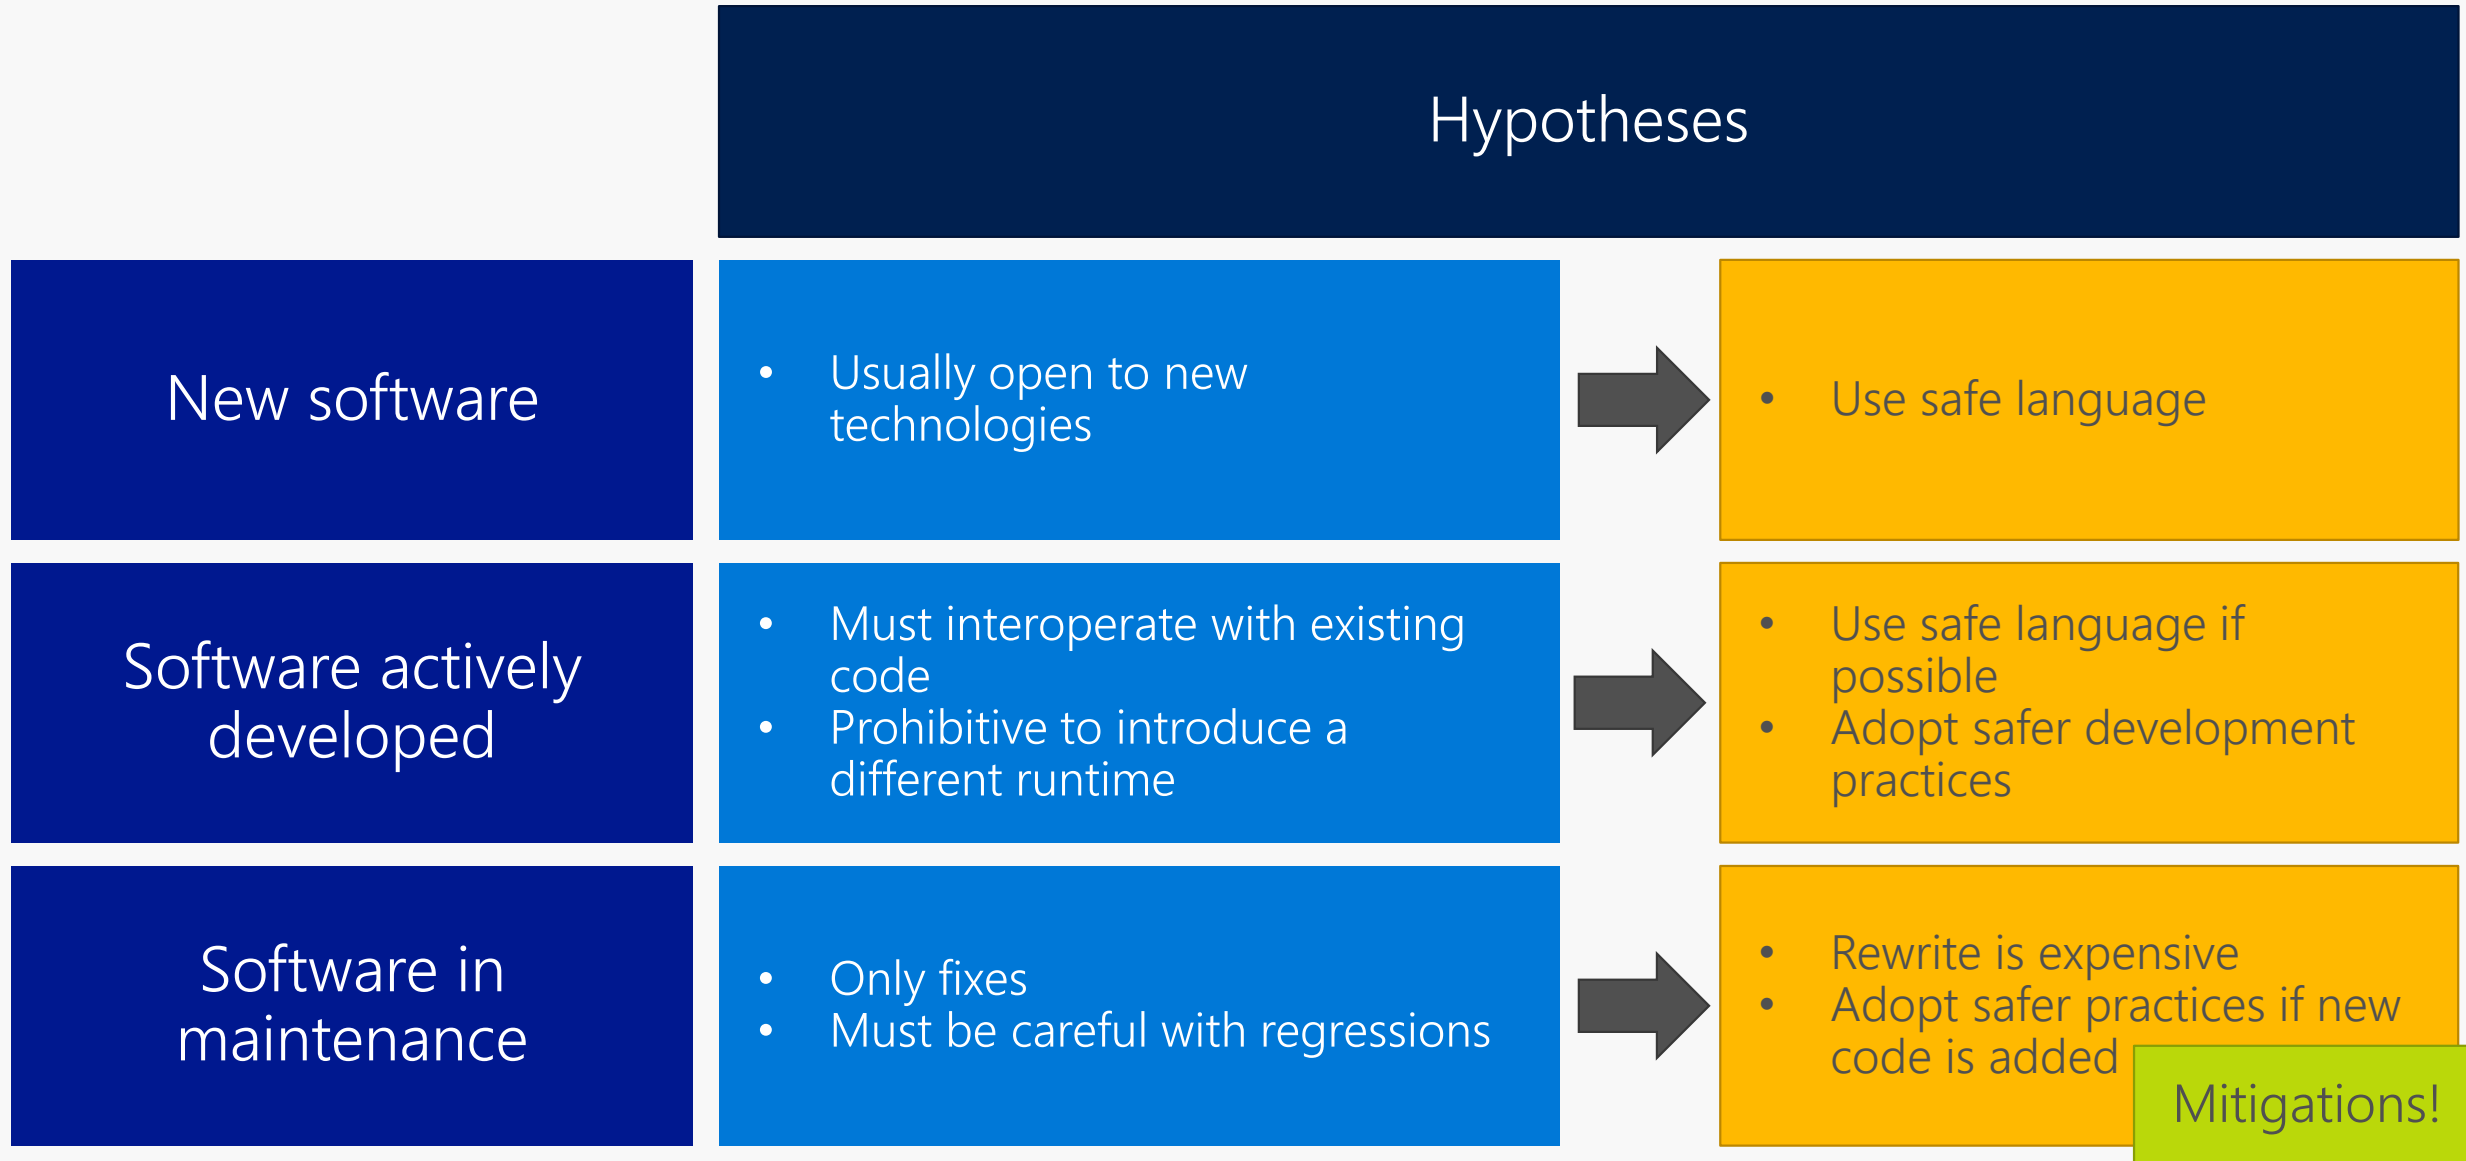
- Rust lifetimes enforced statically
- C++ *gsl::owner<T*>* (and *std::unique_ptr<T>*)

Shared ownership (Reference counted pointers)

- Rust *std::Rc<T>/std::Arc<T>*
- C++ *std::shared_ptr<T>/std::weak_ptr<T>*

Let's move to safer languages!

The stages of software



Conclusion

If possible, switch to safer languages like Rust or .NET

Transition to C++ 17 and enforce the Core Guidelines

In 20 years we shouldn't be introducing any memory safety vulnerabilities 😊



Report vulnerabilities &
mitigation bypasses via our
bounty programs!

<https://aka.ms/bugbounty>

